

Generative Adversarial Networks (GANs)

- Basics of GAN
- Training
- Cost Function Derivative
- Drawbacks
- Implementation in TensorFlow

Generative Adversarial Network (GAN)

- Jun 10, 2014 Ian Goodfellow

Basics of GAN

- Adversarial \rightarrow conflict / opposition
- 2 neural nets competing against each other.
(discriminator, generator)


Objective: GANs are neural networks that are trained in an adversarial manner to generate data mimicking some distribution.

Discriminative Model

- Discriminates between 2 different classes of data (fake / not fake clf).

Generative Model

- A generative model G is trained on a sample data x sampled from true distribution \mathcal{D} is the one which given a random distribution z produces \mathcal{D} which is close to \mathcal{D} in some metric.
 L_1, L_2 norm

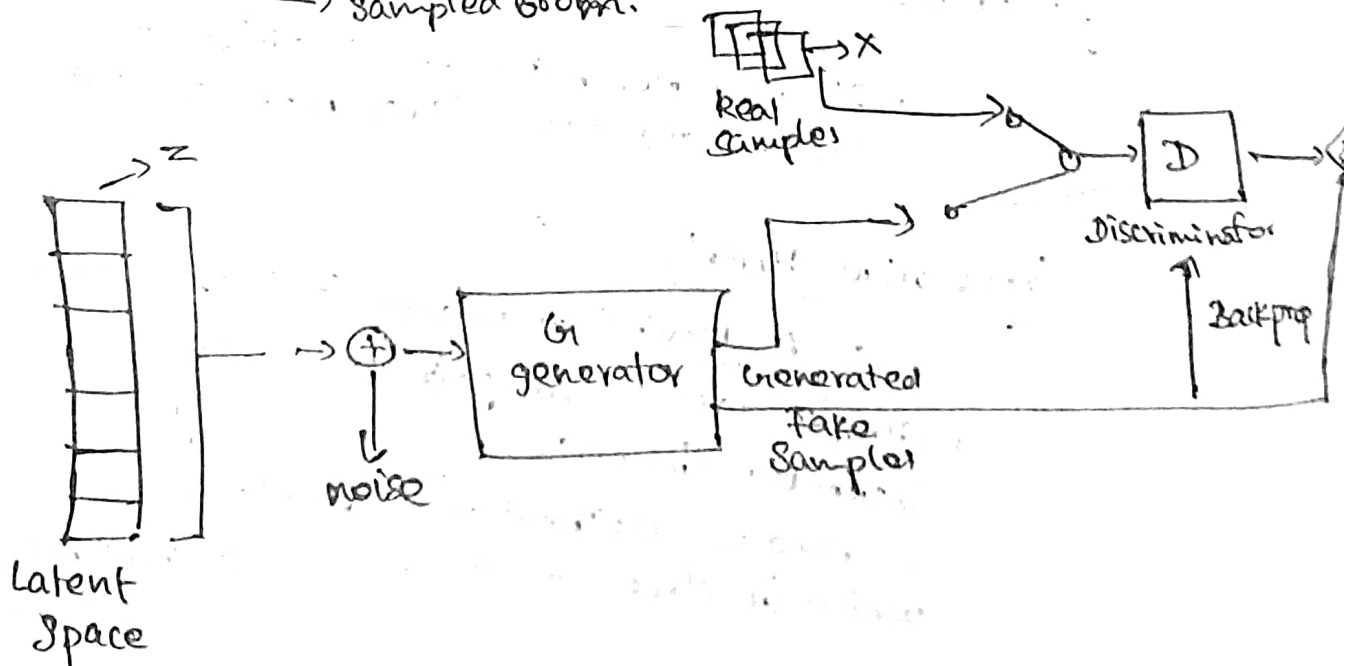


? what is it really doing.

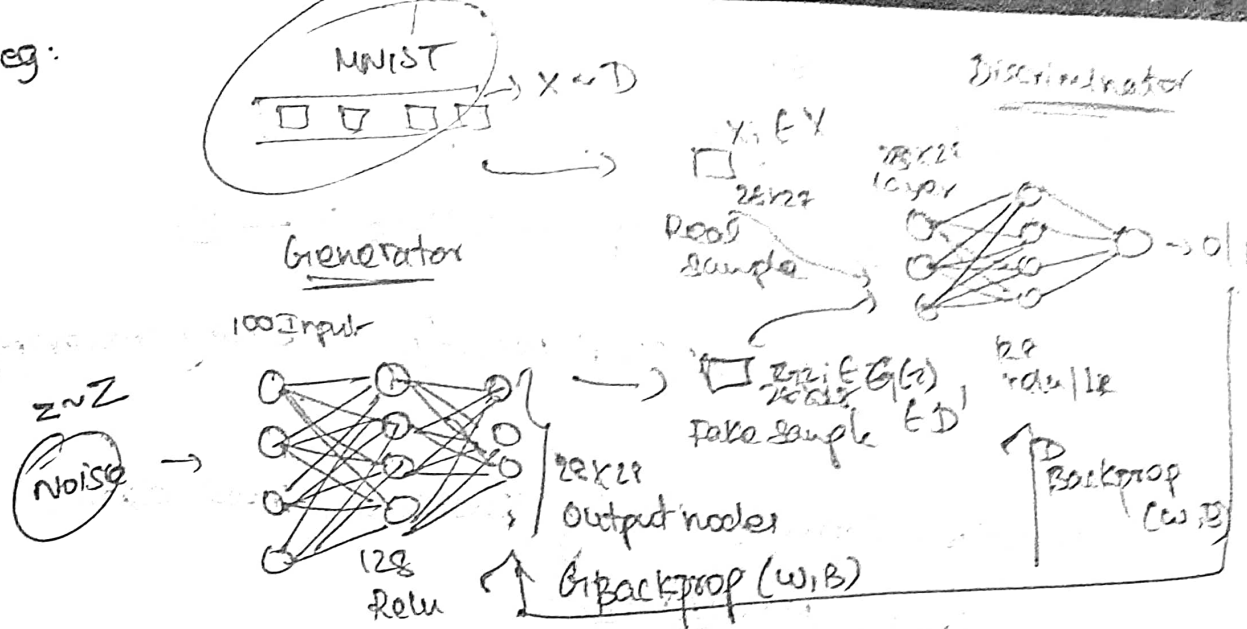
- Take a random distribution z' from Z (true dist)
 - and try converting z' to D' which mimics D (x 's distribution)
 - now we have generated similar-but not same data resembling x 's distribution D .
- $G(z) \longleftrightarrow D'$

$z \sim Z$ maps $G(z) \sim D'$ [$D' \approx D$].

sampled from.



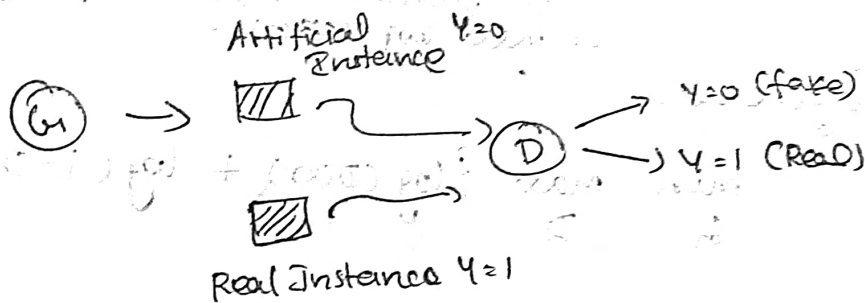
eg:



Generative Adversarial network


Learning mechanism (w, B)

Training the discriminator - D Backprop



D is a simple classifier only. The weights are updated like a classifier neural network.

Training the Generator



Artificial Instance is explicitly named $y=1$ for the D to think it's real.

Then backprop is done to generator based on discriminator output.

Convergence condition: discriminator outputs 0.5

Discriminator

$$L(D(x)) = \log(D(x)) \quad \text{--- (A)}$$

$$L(D(G(z)), \theta) = \log(1 - D(G(z))) \quad \text{--- (B)}$$

$$L = y \log y + (1-y) \log(1-y) \quad \text{--- Binary cross entropy}$$

Objective: correctly classify generated, real data

$$\text{i.e.} \quad \max(A, B)$$

$$\max \{ \log(D(x)) + \log(1 - D(G(z))) \} \quad \text{--- (D)}$$

Generator

$$\min \{ \log(D(x)) + \log(1 - D(G(z))) \} \quad \text{--- (E)}$$

$$\Rightarrow D(G(z)) = 1 \quad (\text{G} \rightarrow \text{give real samples}) \\ \text{or } \max \log(D(G(z)))$$

In general,

$$\min_G \max_D \{ \log(D(x)) + \log(1 - D(G(z))) \}$$

Gradients

$$\nabla_{G_D} = \frac{\partial}{\partial \theta} (\log(D(x)) + \log(1 - D(G(z))))$$

$$\nabla_{G_G} = \frac{\partial}{\partial \theta} \log(1 - D(G(z)))$$

compute gradients based θ which are
parameters of respective neural net
architecture.

Limitations

- Vanishing Gradients.

- use wasserstein loss

- Mode Collapse.

- Generator is set to fool the discriminator and repeatedly produces same outputs.

- ↳ Failure of convergence

- add extra noise to disc. ips.
- penalizing disc. wrt.